

Residual-Correlation Causal Screening with Localized Supervised Modeling: A Practical Framework for Confounder Identification, Error Reduction, and Drift Early-Warning

Rahul Kumar Mandal

Working Paper v1, July 2026 — reporting research conducted June 2022 – present

Abstract. Supervised models routinely fail in deployment because they lean on confounded predictors: variables whose apparent association with the response is inherited from other variables rather than earned through their own causal contribution. We present a practical, fully statistical framework that (i) screens candidate predictors for confounding using an iterative residual-correlation procedure, (ii) ranks the surviving predictors by an empirical *causal power*, (iii) exploits the ranking through causal-factor-based data partitioning with localized, inference-time model training governed by an explicit strength–significance trade-off rule, and (iv) converts the *unaccounted* causal power into an early-warning signal for data and model drift. The screening statistic is the classical partial correlation obtained by the Frisch–Waugh–Lovell decomposition; our contribution is not the statistic but the pipeline built around it: the peeling schedule that prevents a discovered factor’s influence from re-entering through its confounded siblings, the partition-and-localize exploitation strategy, and the missing-causal-power alarm. In commercial deployments spanning credit risk, customer segmentation, and demand forecasting, the framework reduced out-of-sample error by 65–85% when layered over Random Forest, OLS, and Prophet baselines. We state our assumptions — causal sufficiency, no reverse causation, stationary cross-sectional data — plainly, position the method against Granger causality, graph-based discovery, and double machine learning, and address the principal criticisms, including the overfitting objection to localized modeling.

1. Introduction

The purpose of most statistical studies is to establish which variables genuinely drive a response. Modern machine-learning practice largely abandons this question: non-parametric learners optimize predictive loss over arbitrary feature combinations and are indifferent to whether a feature’s apparent contribution is its own or borrowed. The consequence is familiar to every practitioner. Models validate cleanly, ship, and then fail when the environment shifts, because the features carrying the model were *confounded proxies* — variables whose correlation with the response flowed through other variables, observed or not. When that flow changes, the proxy breaks, and no amount of within-distribution validation would have revealed the fragility.

A confounded predictor is a pseudo-explanatory variable: it appears to justify the response, but its explanatory content is largely *interaction effect* — influence inherited from other explanatory variables — rather than *independent effect*. Every explanatory variable can be decomposed conceptually as its own independent effect plus the effect of its interactions with the rest of the universe of variables. When the interaction component dominates, the variable’s importance in any predictive

model is an illusion, and the model built on it carries an uncontrolled failure mode. A mild consequence is inflated error; the extreme consequence is silent model failure under drift, with the data scientist reduced to post-hoc consultation of domain experts whose intuitions cannot be quantified or reproduced.

This paper describes a framework we have developed and deployed commercially over four years to address the confounding component of this problem with purely statistical machinery. We make three practical claims. First, an iterative residual-correlation screen — each round isolating the component of a candidate predictor that is orthogonal to all other candidates, and testing what that component explains of the similarly orthogonalized response — separates variables with genuine independent effect from confounded proxies, and yields a quantitative *causal power* for each. Second, the resulting ranking is directly exploitable: partitioning the training data on the top-ranked factors and training small, local models per partition (only at inference time, and only in the partition containing the query point) reduces error dramatically without the overfitting cost usually associated with data slicing, provided partitioning is stopped by an explicit product rule balancing model strength against statistical significance. Third, the fraction of response variation that the discovered causal factors *cannot* account for — the missing causal power — is a deployable early-warning statistic: when it moves, the causal composition of the environment is changing, and the model is at risk before its accuracy visibly degrades.

We are explicit about what this framework is not. It is not a solution to causal discovery in full generality; no single method is. Granger causality addresses temporal predictive precedence in panel data and is, as is well known, not causality in the interventional sense. Graph- and tree-based discovery methods recover structure under strong assumptions. Randomized experiments and placebo tests address identification but are frequently unavailable. Double machine learning targets unbiased effect estimation for a designated treatment under unconfoundedness. Each solves a part of the puzzle; ours addresses the confounding part in observational, stationary, cross-sectional settings, under assumptions stated in Section 3.

2. Related Work and Positioning

Partial correlation and Frisch–Waugh–Lovell. Our screening statistic is the sample partial correlation of a predictor with the response given all other predictors, computed through the two-residual construction of the Frisch–Waugh–Lovell (FWL) theorem: the coefficient (and correlation) obtained after regressing both the response and the focal predictor on the remaining predictors and correlating the residuals. We claim no novelty for this statistic; it is a century old. The framework’s contribution begins with what is done *around* it: the peeling schedule (Section 4.3), the exploitation strategy (Section 5), and the drift alarm (Section 5.3).

Double machine learning. Robinson’s partialling-out estimator and its modern descendants in double/debiased machine learning (DML) also proceed by residual-on-residual constructions, using flexible learners for the nuisance regressions and cross-fitting for valid inference on a *pre-designated* treatment variable. Our setting differs in aim: we have no designated treatment; we screen and rank *all* candidate predictors, and we consume the ranking not for effect inference but for model construction (partitioning, selection) and monitoring. DML answers “what is the effect of T on Y?”; our pipeline answers “which of my features deserve to be in the model at all, in what order, and how do I exploit and monitor them?”

Constraint-based causal discovery. The PC family of algorithms screens edges by conditional-

independence tests, commonly implemented as partial-correlation tests under Gaussianity. Our procedure can be read as a deliberately restricted, practitioner-oriented relative: rather than searching over graph space, we fix the bipartite question (each predictor versus the response, conditioning on the rest) and add an iterative removal schedule aimed at a specific failure mode — a genuinely causal factor’s influence re-entering the screen through predictors it confounds.

Stepwise selection. Forward stepwise regression also selects variables by incremental explanatory power. Two differences matter. First, our selection statistic is the full-conditioning partial correlation rather than incremental R^2 against the currently selected set, which changes what “spurious” means at each round. Second, and more importantly, selection is only the first third of our pipeline; the partition-and-localize stage and the drift alarm have no stepwise analogue.

Local models and mixtures. Training specialized models on data subsets is an old idea (mixtures of experts, model trees, lazy/local learning). Our variant is distinguished by *what defines the partition* — discovered causal factors in causal-power order, not arbitrary splits optimized for fit — by the explicit strength–significance stopping rule, and by inference-time localization: exactly one local model is trained, in the single partition containing the query point.

3. Assumptions and Problem Setup

Let Y denote the response and $X = (X_1, \dots, X_p)$ the candidate explanatory variables, observed as n i.i.d. cross-sectional rows. We work under four explicit assumptions.

A1 (Stationarity, cross-sectional design). The data are stationary and cross-sectional; temporal dynamics are out of scope. (The drift alarm of Section 5.3 monitors *departures* from this assumption in deployment.)

A2 (Causal sufficiency). All relevant variables of the universe under study are included in X . In practice this is a simplifying idealization limited by the researcher’s knowledge and data availability; Section 7 discusses the consequences of its failure.

A3 (No reverse causation). Y does not cause any X_j ; direction is from explanatory variables to response. Combined with the cross-sectional design, this also excludes feedback loops.

A4 (Temporal precedence, encoded by design). Explanatory events precede the response event. In cross-sectional data this is a design responsibility — features must be measured strictly before the response window.

We also note the classical completeness condition on true causality: if a set of causal factors fully explains the response, no additional true causal factor of that response can exist in the universe. This motivates our use of *unexplained* residual variation as “missing causal power” (Section 5.3) rather than as noise.

Under A1–A4, the target of the screening stage is: for each j , quantify the component of association between X_j and Y that is *not* attributable to the remaining variables X_{-j} .




Figure 1: **Figure 1.** The framework pipeline. Training data enters the residual-correlation screen with iterative peeling (Sec. 4), yielding a causal-power ranking. The ranking drives data partitioning under the strength×significance stopping rule (Sec. 5.1); at inference time, in-support queries receive one locally trained model (interpolation) while out-of-support queries fall back to the global model (extrapolation). The unexplained residual — missing causal power — is monitored as a drift early-warning statistic (Sec. 5.3).

4. The Screening Procedure

4.1 The two-model residual construction

Fix a focal variable, say X_1 , and pose the working hypothesis H_0 : X_1 is a true causal factor of Y (against the alternative that it is not). Estimate two auxiliary regressions:

- **Model 1 (response side):** regress Y on $X_{-1} = (X_2, \dots, X_p)$. The residual $e_Y = Y - \widehat{E}[Y | X_{-1}]$ contains, by construction, all information about Y not explained by the other variables.
- **Model 2 (predictor side):** regress X_1 on X_{-1} . The residual $e_1 = X_1 - \widehat{E}[X_1 | X_{-1}]$ contains all information about X_1 not explained by the other variables.

The screening statistic is the correlation of the residuals,

$$\hat{\rho}_1 = \text{corr}(e_Y, e_1),$$

which we term the **causal power** of X_1 (relative to the universe X , under A1–A4). The logic is immediate: if H_0 holds, the part of Y unexplained by the others (e_Y) is precisely the contribution of X_1 , and the part of X_1 unexplained by the others (e_1) is its independent content; the two must align, ideally perfectly ($\hat{\rho}_1 \rightarrow 1$) if X_1 were the *only* causal factor. If instead X_1 ’s apparent association with Y is inherited — inflated or deflated by the other variables — then stripping the others removes it, and $\hat{\rho}_1$ collapses toward zero. Statistically, $\hat{\rho}_j$ is the sample partial correlation $\rho_{YX_j \cdot X_{-j}}$; its significance can be assessed by the standard t -test or Fisher z -transform at sample size n and conditioning dimension $p - 1$.

The auxiliary regressions are linear in our deployments (OLS), which keeps the procedure transparent, fast, and exactly FWL-consistent; flexible learners can be substituted at the cost of interpretability, moving the construction toward the DML nuisance style.

4.2 Interpretation

Three readings of $\hat{\rho}_j$ are used operationally. *Magnitude* ranks variables by independent contribution. *Sign* is diagnostic: a variable whose raw correlation with Y is large but whose causal power is near zero is a confounded proxy — exactly the fragile feature that predictive validation cannot flag. *Instability* of $\hat{\rho}_j$ across resamples signals borderline factors to be handled conservatively.

4.3 Iterative peeling

A single pass is insufficient, for a specific reason: a genuinely causal factor influences not only Y but typically also other candidate predictors. If it remains in the conditioning set of later tests, its effect re-enters the screen *through* the variables it confounds, distorting their apparent powers. The procedure therefore *peels*:

1. Compute $\hat{\rho}_j$ for all remaining candidates; let $j^* = \arg \max_j |\hat{\rho}_j|$ (subject to significance).
2. Record X_{j^*} and its causal power; **remove it from the candidate set and from all subsequent conditioning sets.**
3. Repeat on the reduced set until no remaining candidate is significant, or a preset budget is reached.

The output is an ordered list $X_{(1)}, X_{(2)}, \dots$ with causal powers $\hat{\rho}_{(1)}, \hat{\rho}_{(2)}, \dots$ — a decreasing-order causal ranking of the universe. Removal at each round is the design choice that distinguishes the schedule from one-shot partial-correlation screening: it prevents double-counting a discovered factor’s influence via its confounded siblings.

Remark (peeling vs. reverse causation). Because discovered factors are removed sequentially, the schedule never revisits the possibility that Y drives an already-selected factor; this is coherent only because A3 excludes reverse causation by assumption. The procedure does not test A3 and inherits its failure modes (Section 7).

5. Exploiting the Ranking

5.1 Causal-factor partitioning with localized inference-time training

Given the ranking, partition the training data on the unique values (categorical/ordinal, or binned continuous/temporal values) of $X_{(1)}$; within each cell, optionally partition further on $X_{(2)}$, and so on. Two opposing forces govern the depth of partitioning: model strength inside a cell (error decreases as cells become causally homogeneous — the aligned causal factors add their strength) and statistical significance (which decays as cell sample size shrinks). We resolve the trade-off with an explicit stopping rule:

Rule (strength $\times$ significance). Partition until the product (model strength) $\times$ (statistical significance) is maximized; when a query point falls *inside* a cell’s support, strength may be weighted more aggressively; when it falls *outside* all cells, do not use the partition at all — fall back to the global model, whose larger sample retains significance and generalization.

Two engineering decisions make this practical. First, **inference-time localization**: no models are pre-trained per cell. When a query arrives, it is routed to its cell and one small model is trained there and then — on a fraction of the data, hence faster than one global training run — and used for that prediction. Second, **interpolation/extrapolation split**: in-cell queries are interpolation

and receive the local model; out-of-support queries are extrapolation and receive the global model. This split is the answer to the overfitting objection (Section 7).

5.2 Synthetic sparse regularization

A second exploitation mode retains only the top- k causal factors and drops the remainder before any supervised modeling. Because the dropped variables are precisely those whose association with Y is dominantly inherited, this acts as a *causally targeted* sparsification — a synthetic regularizer that removes the features most likely to break under drift, unlike norm-based penalties which are agnostic to why a feature predicts.

5.3 Missing causal power as a drift alarm

By the completeness condition of Section 3, if the discovered factors were the whole causal story, the ranked powers would account for the response entirely. The shortfall — the residual variation of Y unexplained after the discovered factors, i.e., the causal power *missing from the dataset* — is itself informative. Deployed, it becomes a monitored statistic with a threshold trigger: a rising missing-power reading means causal content is leaving the observed universe (data drift) or the response’s causal composition is changing (model drift), and it moves *before* headline accuracy collapses, because it is computed on the causal decomposition rather than the loss. This converts the framework from a training-time tool into a lifecycle guardrail: a preventive alarm available during the entire deployment, rather than a post-mortem explanation after failure.

6. Deployment Results

The framework has been applied over four years across commercial problems including credit risk, customer segmentation, demand forecasting, and econometric analysis, layered over heterogeneous base learners.

- **Credit risk (global automotive manufacturer).** A causal-statistical credit-risk model built with the screening and selection stages preceded a 42% reduction in payment defaults in the quarter following deployment.
- **Model-enhancement deployments.** Used as a pre-modeling layer over a Random Forest classifier and over OLS baselines at multinational clients, and over Facebook’s Prophet forecaster, the framework reduced out-of-sample error by 65–85% across engagements (85% at the upper end for high-dimensional prediction with confounder removal; 65% for the Prophet enhancement).

We report these as *case results from production deployments*, with the important caveat that client data cannot be released. A reproducible public-benchmark suite (synthetic ground-truth confounding, plus standard public datasets, with DoWhy/EconML baselines) is in preparation as the companion open-source release of this paper, and is the appropriate venue for like-for-like comparison.

7. Criticisms and Responses

“Partitioning then fitting per cell is overfitting.” Overfitting is a property of *one* model asked to both memorize and generalize. Here the two jobs are held by different models: many small local models handle in-support queries (interpolation — where causal homogeneity, not memorization,

drives the low error), and the single global model handles out-of-support queries (extrapolation — where its larger sample preserves significance). Because no local model is ever asked to generalize beyond its cell, the classical overfitting failure — confident nonsense on unseen regions — is structurally excluded rather than merely penalized.

“Per-cell training is computationally expensive.” It would be if all cells were trained eagerly. Under inference-time localization exactly one local model is trained per query, on a cell that is by construction much smaller than the full dataset; local training is therefore *cheaper* than one global refit, and cells can be cached.

“Removing confounders does not prove causality.” Correct, and we do not claim it. The framework addresses the confounding component of causal analysis under A1–A4. If reverse causation exists in reality, A3 is violated and the peeling schedule will misattribute direction; if the researcher’s universe omits a true factor (A2 violated — the common case, since inclusion is limited by the researcher’s knowledge, bias, and data availability at design time), discovered powers are relative to the observed universe and the omission surfaces only as missing causal power, not as an identified variable. These are limitations of scope shared, in various combinations, by every observational method; we prefer stating them to overclaiming. What the framework *does* deliver under its assumptions is reproducible, expert-free identification of which observed variables carry independent explanatory content, with a quantified account of what remains unexplained — transparent and reproducible at any analyst’s end, in contrast to domain-expert adjudication.

8. Limitations and Future Work

Beyond A1–A4: (i) linear auxiliary regressions bound the screen to linear independent effects — nonlinear nuisance learners are a natural extension, at the cost of FWL exactness; (ii) significance handling across many peeled tests warrants a formal multiple-testing treatment; (iii) the strength×significance rule is currently tuned empirically, and a theoretical characterization of its optimum is open; (iv) extension to temporal and panel data — where A1 and A4 interact with dynamics — is the most requested direction from deployments; and (v) the public benchmark suite (synthetic confounding with known ground truth; comparisons against DoWhy/EconML pipelines, stepwise selection, and DML variable-importance baselines) is in preparation.

9. Conclusion

Confounded predictors are the quiet cause of a large share of production model failures. We described a deliberately simple, fully statistical pipeline — residual-correlation screening via the classical two-regression construction, an iterative peeling schedule, causal-power ranking, partition-based localized modeling with an explicit strength–significance stopping rule, and a missing-causal-power drift alarm — that has repeatedly cut out-of-sample error by 65–85% in commercial deployments and provides a preventive, quantitative alternative to domain-expert post-mortems. The statistic is classical; the pipeline, its stopping rule, its inference-time localization, and its use of unexplained causal power as a lifecycle alarm are, to our knowledge, a practically distinct combination. We state our assumptions plainly and invite scrutiny through the forthcoming open-source benchmark release.

References

(to be completed in the *LaTeX/submission pass*)

- Frisch, R., Waugh, F. V. (1933). Partial Time Regressions as Compared with Individual Trends. *Econometrica* 1(4).
- Lovell, M. C. (1963). Seasonal Adjustment of Economic Time Series and Multiple Regression Analysis. *JASA* 58.
- Robinson, P. M. (1988). Root-N-Consistent Semiparametric Regression. *Econometrica* 56(4).
- Chernozhukov, V., et al. (2018). Double/Debiased Machine Learning for Treatment and Structural Parameters. *Econometrics Journal* 21(1).
- Spirtes, P., Glymour, C., Scheines, R. (2000). *Causation, Prediction, and Search*. MIT Press.
- Granger, C. W. J. (1969). Investigating Causal Relations by Econometric Models and Cross-Spectral Methods. *Econometrica* 37(3).
- Pearl, J. (2009). *Causality: Models, Reasoning, and Inference*. 2nd ed., Cambridge University Press.
- Jacobs, R. A., Jordan, M. I., et al. (1991). Adaptive Mixtures of Local Experts. *Neural Computation* 3(1).
- Sharma, A., Kiciman, E. (2020). DoWhy: An End-to-End Library for Causal Inference. arXiv:2011.04216.
- Taylor, S. J., Letham, B. (2018). Forecasting at Scale (Prophet). *The American Statistician* 72(1).